

Steering a Network with Its Own Reflection: True, False, and Noise Self-Model Broadcasts in an Evolving Multi-Agent Network

Ivan Bhargava

Independent Researcher

`ivan.bhargava@example.com`

WWW home page: <https://github.com/Ivan825/Project-ONE>

Abstract. Many networks respond to measurements of their own state: markets to indices, organizations to metrics, platforms to trending signals. We present a controlled laboratory for this reflexive loop: an evolving agent population forms an adaptive network whose macrostate is compressed into a self-model and broadcast back — accurately, distorted, or as matched-bandwidth noise. Across 1,092 deterministic runs with paired seeds and pre-registered outcomes: measurement without broadcast is causally inert; a false self-model pulls the macrostate toward itself far more strongly than noise (Cliff’s δ up to 0.85); which lies come true depends on the response regime — an alarmist lie that is self-defeating under corrective responses becomes strongly self-fulfilling under conformist ones, while a flattering lie flips from inert to benevolent; and heritable attention to the broadcast is selected against in proportion to its unreliability: populations fed false self-models evolve to ignore their observer.

Keywords: adaptive networks, agent-based modeling, reflexivity, self-fulfilling dynamics, evolution of trust, network resilience

1 Introduction

Adaptive networks — graphs whose topology coevolves with the dynamical states of their nodes — are a central object of network science [1, 2]. A distinctive subclass has received far less controlled attention: networks that respond to *descriptions of themselves*. Financial markets move on published indices computed from their own trades [3]; universities restructure around rankings derived from their own behavior [4]; recommendation platforms amplify items *because* an internal measurement declared them trending. In each case a macroscopic measurement of the system is compressed, published, and re-enters the system as an input to microscopic decisions — a loop Merton called the self-fulfilling prophecy [5] and whose measure-distorting variant is known as Goodhart’s law [6].

Empirical study of this loop is confounded: in real systems one can rarely withhold the measurement, falsify it in controlled directions, or replay history under identical randomness. We therefore built a minimal simulated world in




Fig. 1. The reflexive loop and its experimental corruption points. Agents’ local actions generate the network; the observer compresses it into a self-model; the broadcast channel returns that self-model truthfully, distorted, as noise, or not at all. Each agent’s response is scaled by its heritable sensitivity γ_i and a global gain g .

which the loop can be opened, closed, and corrupted at will, and where every run is exactly reproducible from a configuration and seed. Agents are deliberately simple — transparent trait vectors rather than learned policies — so that every behavioral pathway from broadcast to response is inspectable.

Our contributions are: (1) an open, deterministic experimental platform coupling an evolving agent ecology, an adaptive interaction network, a global observer, and a controllable broadcast channel; (2) a five-condition causal design that separates the effect of *being measured* from the effect of *being told*, and the effect of *self-information* from that of *stimulation* (via a matched-bandwidth noise control); (3) evidence, robust across feedback gains and two shock severities, that false self-models steer the macrostate toward themselves while noise does not, and that the steering direction is a function of the distortion type; and (4) to our knowledge the first demonstration in an ecological agent-based setting that *attention to a collective self-model is itself an evolvable trait*, selected against in proportion to the self-model’s unreliability. Figure 1 summarizes the design.

2 Related Work

Adaptive and temporal networks. Coevolution of state and topology is surveyed in [1, 2]; our model belongs to this family, with births, deaths and rewiring driven by agent decisions. Robustness of networks to random and targeted node removal is classical [7]; we use targeted hub removal as a standardized perturbation rather than an object of study. **Global coupling in physics.** Globally coupled oscillators and maps show that a mean field fed back to units reshapes collective dynamics [8]. Our design differs in two ways central to the research question: the “mean field” is a multidimensional, potentially *false* description, and units are born, die, and evolve. **Reflexivity and reactive measurement.** Self-fulfilling prophecy [5], reflexivity in markets [3], and the reactivity of rankings and metrics [4, 6] motivate our false-broadcast conditions; we contribute a platform where these effects are dosable and ablatable. **Signaling and the evolution of trust.** In sender–receiver games, receivers can evolve to ignore unreliable signals [9]. Our distrust result is consonant with that theory but arises

in an open ecological population where the signal is *self-referential* — a compressed description of the receiving population itself — and where attention to it is a continuous heritable trait under selection from an energy economy, not a payoff matrix. **Agent-based modeling.** Agent-based studies of collective opinion and cooperation are a mainstay of this community [16, 17]; our methodology follows the generative ABM tradition [10, 11]; the engine uses ideas from Mesa [12] and NetworkX [13]. Multi-agent reinforcement learning studies global state as an *architectural* aid [14]; we instead treat global information as an *experimental variable* with controlled veracity. Populations of language-model agents are now deployed with shared dashboards describing their own collective state [15]; the present rule-based setting provides causally clean baselines for that emerging regime.

3 The Model

3.1 Agents and ecology

Each agent i carries a heritable trait vector $\theta_i = (\text{coop}, \text{expl}, \text{risk}, \text{soc}, \text{share}, \gamma_i) \in [0, 1]^6$, where γ_i is its *global sensitivity* — how strongly broadcasts modulate its behavior. Agents hold energy; each step they pay metabolism and choose one action — harvest, share (costly transfer to a neighbor), connect, prune, reproduce, or idle — sampled with probabilities proportional to trait- and state-dependent propensities. The resource pool regenerates logistically, giving density-dependent competition. Reproduction requires an energy threshold and copies traits with Gaussian mutation ($\sigma = 0.05$); energy exhaustion or age causes death. Founding ages are staggered to avoid cohort artifacts.

3.2 Observer and broadcast conditions

Every $\Delta t = 10$ steps the observer computes the macrostate $S(t)$: population, fragmentation (1 minus largest-component share), degree- and betweenness-concentration, Freeman centralization [19], per-capita costly-helping rate, energy Gini, and turnover, plus mean trait values (logged, never broadcast). Five conditions define what agents receive (Table 1). Distorted broadcasts come in three modes: *invert* ($b_k = 1 - s_k$ per normalized metric), *crisis* (fixed alarming values), and *utopia* (fixed reassuring values). Broadcast values shift action propensities through fixed corrective response rules — e.g., reported fragmentation above 0.3 raises the propensity to form links; a reported cooperation deficit raises the propensity to share — each scaled by γ_i and a global gain g .

3.3 Emergent network structure

The interaction networks that emerge are non-trivial: at $t = 2,000$ the giant component exhibits clustering $1.8\text{--}2.4\times$ an Erdős-Rényi null with matched size and density, at comparable characteristic path lengths ($L \approx 1.9\text{--}2.5$) — a small-world-like regime — with a broad degree spread (coefficient of variation $0.6\text{--}0.75$;

Table 1. Experimental conditions. All start from matched parameter distributions with paired random seeds.

Condition	Observer Agents receive		Role
A — Local only	off	nothing	baseline
B — Observed, blind	on	nothing	controls for measurement
C — True feedback	on	accurate $S(t)$	main treatment
F — False feedback	on	distorted $S(t)$	reflexivity probe
N — Noise feedback	on	uniform random, same keys/ranges	stimulation control

maximum degree $2\text{--}2.5\times$ the mean). The substrate on which feedback acts is therefore a heterogeneous, clustered network rather than a homogeneous mixing population.

3.4 Reproducibility

The simulation is deterministic given (configuration, seed): one PRNG drives all choices, agents act in sorted order, and a state hash certifies replay. Identical hashes were verified across operating systems and machines. All code, configurations, and analysis scripts are public; every figure regenerates from two commands.

4 Experimental Design

Four outcomes were fixed before the first campaign: (1) post-shock recovery time (steps until population regains 90% of its pre-shock mean after standardized hub removal); (2) mean post-shock fragmentation; (3) per-capita costly-helping rate; (4) *story pull* — the mean signed movement of the macrostate toward the broadcast, counted only where broadcast and reality differ by more than 0.02, so that the measure is well defined under false and noise broadcasts. Campaigns: a *flagship* run (conditions A/B/C/ F_{invert} /N $\times$ 50 paired seeds, mild shock), a *harsh-shock* replication (removal of 40% of hubs, $\times 30$ seeds), and a *sensitivity sweep* (feedback gain $g \in \{0.2, 0.8, 1.6\} \times \{A, C, F_{\text{invert}}, F_{\text{crisis}}, F_{\text{utopia}}, N\} \times 20$ seeds), a *scale check* re-running the core conditions with doubled population and resource capacity, and a *response-regime campaign* repeating all broadcast conditions with conformist agents (15 seeds per cell): 1,092 runs in total. Group differences use Mann–Whitney U with Cliff’s δ effect sizes; all other logged measures are exploratory.

5 Results

Table 2 summarizes the pre-registered comparisons; the following subsections unpack them.

Table 2. Primary comparisons across campaigns (medians; Mann–Whitney U ; Cliff’s δ). FL = flagship (50 seeds/condition), HS = harsh shock (30 seeds/condition).

Outcome	Comparison	Medians	δ	p
every outcome (FL, HS)	B vs A	identical	0.000	1.0
story pull (FL)	F vs N	0.0067 / 0.0018	0.76	6×10^{-11}
story pull (HS)	F vs N	0.0078 / 0.0021	0.85	1.6×10^{-8}
cooperation (FL)	C vs A	0.397 / 0.179	0.96	1.5×10^{-16}
cooperation (HS)	C vs A	0.399 / 0.176	0.95	2.6×10^{-10}
fragmentation (FL)	F vs A	0.021 / 0.042	-0.50	2.0×10^{-5}
$\Delta\bar{\gamma}$ (HS)	F vs A	-0.42 / -0.03	-0.94	4.2×10^{-10}
$\Delta\bar{\gamma}$ (HS)	N vs A	-0.30 / -0.03	-0.77	2.8×10^{-7}
recovery time (FL, HS)	C, F, N vs A	$\approx$ equal	$ \delta < 0.19$	> 0.12

5.1 Measurement alone is causally inert

Under paired seeds, condition B is trajectory-identical to A: $\delta = 0.000$ on every outcome in both the flagship and harsh-shock campaigns. The instrument acquires causal power only when it broadcasts — a validation that the design isolates the channel of interest.

5.2 A false self-model pulls the macrostate toward itself

Story pull under the inverted false broadcast exceeds the noise control at every gain tested (flagship: median 0.0067 vs 0.0018, $\delta = 0.76$, $p \approx 6 \times 10^{-11}$; harsh shock: 0.0078 vs 0.0021, $\delta = 0.85$, $p \approx 1.6 \times 10^{-8}$). Relative to noise the effect is strongest at *low* gain (ratio ≈ 7 at $g = 0.2$): a quiet, systematic lie out-steers loud noise, while at high gain noise itself begins to steer (Fig. 2, middle). Spillovers are substantial: the inverted broadcast roughly halves post-shock fragmentation relative to baseline ($\delta = -0.50$, $p \approx 2 \times 10^{-5}$) because agents persistently told the network is fragmenting over-connect it (Fig. 3).

5.3 The type of lie sets the direction

The three distortion modes produce qualitatively different regimes (Fig. 2). The *inverted* self-model, always reporting the mirror of reality, generates the strongest pull: it keeps every corrective response permanently active and thereby parasitizes the population’s repair machinery. The *crisis* self-model is self-defeating for structure — corrections push fragmentation away from the alarmist report — yet its permanent cooperation-deficit signal mobilizes the highest costly-helping rates observed (0.36–0.45). The *utopia* self-model is behaviorally inert: because responses fire only on reported deficits, a flattering broadcast triggers nothing and is indistinguishable from disconnecting the channel, at every gain. Under purely corrective response rules, flattery equals silence.




Fig. 2. Sensitivity sweep (40% hub-removal shock; 20 seeds per cell; medians). Left: evolved change in mean global sensitivity γ — attention to the self-model is selected against in proportion to broadcast unreliability and gain; the flattering (utopia) lie tracks the no-feedback baseline exactly. Middle: story pull toward the broadcast. Right: cooperation rate.

Table 3. The same lie under two response regimes (medians; harsh shock, $g = 0.8$). Baseline (A): cooperation 0.18, fragmentation 0.038.

Lie	Corrective		Conformist	
	coop.	fragm.	coop.	fragm.
Invert	0.27	0.021	0.42	0.054
Crisis	0.36	0.025	0.26	0.400
Utopia	0.19 (inert)	0.037 (inert)	0.428	0.025

5.4 The response regime decides which lies come true

Repeating all broadcast conditions with conformist agents reverses the taxonomy exactly as the mechanism predicts (Table 3). The *crisis* lie, self-defeating under correction, becomes strongly self-fulfilling under conformity: told the world is fragmenting, imitators cut ties and fragmentation rises an order of magnitude above baseline (0.400 vs 0.038). The *utopia* lie flips from behaviorally inert to benevolently self-fulfilling: “everyone cooperates” becomes true by imitation (cooperation 0.428, the highest observed in any condition). Under conformity noise steers as strongly as the inverted lie, since imitation amplifies any reported value. Which false self-models come true is therefore not a property of the lie alone but of the lie–response pair — the reflexive analogue of the observation that a measure’s distorting power depends on how actors respond to it.

5.5 Populations fed lies evolve to stop listening

Mean global sensitivity $\bar{\gamma}$ changes over ~ 40 generations as an emergent, unprogrammed function of broadcast quality (Fig. 4). At $g = 0.8$ (harsh-shock campaign): no broadcast -0.03 (neutral drift), truth -0.10 , noise -0.30 ($\delta = -0.77$), false -0.42 ($\delta = -0.94$, $p \approx 4 \times 10^{-10}$). The sweep reveals a dose–response: distrust of the false broadcast deepens from -0.11 ($g = 0.2$) to -0.45 ($g = 1.6$)

The same world under two mirrors: identical seed, $t = 2000$

A — no feedback
mean degree 11.2 · cooperation 0.11 · fragmentation 0.000

F — false feedback (invert)
mean degree 22.4 · cooperation 0.21 · fragmentation 0.000




Fig. 3. The same world under two mirrors. Two runs from an identical seed at $t = 2000$: without feedback (left) and under the inverted false broadcast (right), which persistently reports fragmentation and thereby drives over-connection (mean degree 22.4 vs 11.2) and elevated costly helping. Node size encodes energy; color encodes generation.

— and at the highest gain even *accurate* feedback is selected against (-0.24): a broadcast that perturbs behavior strongly becomes a fitness liability even when true. The utopia lie, being consequence-free, breeds no distrust at all.

5.6 Accurate feedback as a distributed homeostat

True feedback roughly doubles costly helping (flagship: 0.397 vs 0.179, $\delta = 0.96$, $p \approx 10^{-16}$; replicated under harsh shock, $\delta = 0.95$), rising monotonically with gain ($0.27 \rightarrow 0.49$); Fig. 5 shows the sustained divergence of trajectories. The loop functions as a population-level thermostat: reported deficits are repaired by distributed individual action.

5.7 A robust negative result

Post-shock recovery time was insensitive to feedback in every condition and campaign, even with 40% of hubs removed (all $p > 0.12$; median ≈ 40 steps). In this ecology, resilience is demographic — fast regrowth — rather than informational. We report this as a boundary of the phenomenon rather than adjusting the protocol until the null disappears.

5.8 Robustness to system scale

Re-running the core conditions (A, C, F_{invert} , N) with doubled initial population and proportionally scaled resources (steady-state population ≈ 120 – 130 vs. ≈ 60) preserves every effect: cooperation under true feedback is 0.399 vs. 0.188




Fig. 4. Evolution of attention to the self-model (harsh-shock campaign, $g = 0.8$, 30 seeds per condition). Change in population-mean global sensitivity $\bar{\gamma}$ between the start and end of the run. The ordering (false < noise < true < none) emerged from selection alone.

baseline (vs. 0.397/0.179 at reference scale), story pull retains $\delta = 0.78$ (F vs. N), and the distrust ordering is intact with F vs. A at $\delta = -1.0$ ($p \approx 1.6 \times 10^{-4}$, $n = 8$ seeds per condition). The phenomena are not finite-size artifacts of the reference configuration.

6 Discussion and Limitations

The first-order *directions* of response are encoded in the agents’ corrective rules; that is by design, and we do not claim them as discoveries. The findings live one level up: the magnitude and asymmetry of steering (consistent lies $\gg$ matched noise, strongest at low gain), the spillovers to non-targeted structure, the qualitative taxonomy of lies (exploitative / mobilizing-but-self-defeating / inert), and above all the evolutionary layer, in which the credibility of a collective self-model becomes a selected property of the population. Both response regimes tested here are fixed rules; learned or heterogeneous response policies are the natural next step. Other limitations: a single ecology and parameter family (though effects are monotone across an $8\times$ gain range and stable under a $2\times$ scale change), rule-based rather than learned agents, and one network per run rather than interacting populations.

The results have a practical reading for engineered systems. Metric-driven infrastructure — autoscalers, observability pipelines, trending algorithms — constitutes exactly this loop; the false-broadcast conditions model telemetry corruption, and the utopia result implies that *flattering telemetry is operationally equivalent to no telemetry* while raising no alarms. Documented episodes of measurement reshaping the measured system — from ranking reactivity in institutions [4] to the degradation of Google Flu Trends once its predictions entered




Fig. 5. Mean cooperation trajectories by condition (flagship campaign, 50 seeds per condition; dashed line marks the standardized shock). Conditions A and B coincide exactly. True and noise feedback sustain elevated cooperation; the inverted false broadcast sustains an intermediate level.

the environment they forecast [18] — are natural empirical targets for the mechanisms isolated here, and connecting the model’s regimes to such datasets is the clearest path to empirical grounding. For multi-agent AI, where fleets of agents increasingly receive dashboards describing their own collective state, the evolution-of-distrust result suggests that persistent misreporting does not merely misdirect a population — it degrades the control channel itself.

7 Conclusion

We introduced a reproducible causal laboratory for reflexive self-modeling in adaptive networks and used it to show that (i) being measured changes nothing while being told changes much; (ii) a consistent false self-description steers a collective toward itself far more effectively than noise; (iii) the kind of lie determines whether it exploits, mobilizes, or pacifies; and (iv) selection acting on a heritable attention trait makes a population’s trust in its own self-model an evolving quantity — one that unreliable observers spend down. Future work extends the platform to conformist response rules, continuity under total population turnover, and learned (including language-model) agents.

Reproducibility. Code, configurations, all 970 run manifests, and one-command reproduction of every figure: github.com/Ivan825/Project-ONE.

References

1. Gross, T., Blasius, B.: Adaptive coevolutionary networks: a review. *J. R. Soc. Interface* 5(20), 259–271 (2008)

2. Holme, P., Saramäki, J.: Temporal networks. *Phys. Rep.* 519(3), 97–125 (2012)
3. Soros, G.: Fallibility, reflexivity, and the human uncertainty principle. *J. Econ. Methodol.* 20(4), 309–329 (2013)
4. Espeland, W.N., Sauder, M.: Rankings and reactivity: how public measures recreate social worlds. *Am. J. Sociol.* 113(1), 1–40 (2007)
5. Merton, R.K.: The self-fulfilling prophecy. *Antioch Rev.* 8(2), 193–210 (1948)
6. Strathern, M.: ‘Improving ratings’: audit in the British university system. *Eur. Rev.* 5(3), 305–321 (1997)
7. Albert, R., Jeong, H., Barabási, A.-L.: Error and attack tolerance of complex networks. *Nature* 406, 378–382 (2000)
8. Kuramoto, Y.: *Chemical Oscillations, Waves, and Turbulence*. Springer (1984)
9. Skyrms, B.: *Signals: Evolution, Learning, and Information*. Oxford University Press (2010)
10. Epstein, J.M., Axtell, R.: *Growing Artificial Societies: Social Science from the Bottom Up*. Brookings/MIT Press (1996)
11. Bonabeau, E.: Agent-based modeling: methods and techniques for simulating human systems. *PNAS* 99(suppl. 3), 7280–7287 (2002)
12. Kazil, J., Masad, D., Crooks, A.: Utilizing Python for agent-based modeling: the Mesa framework. In: *SBP-BRiMS*. Springer (2020)
13. Hagberg, A., Schult, D., Swart, P.: Exploring network structure, dynamics, and function using NetworkX. In: *SciPy* (2008)
14. Foerster, J., Assael, Y.M., de Freitas, N., Whiteson, S.: Learning to communicate with deep multi-agent reinforcement learning. In: *NeurIPS* (2016)
15. Park, J.S., O’Brien, J., Cai, C.J., Morris, M.R., Liang, P., Bernstein, M.S.: Generative agents: interactive simulacra of human behavior. In: *UIST* (2023)
16. Castellano, C., Fortunato, S., Loreto, V.: Statistical physics of social dynamics. *Rev. Mod. Phys.* 81(2), 591–646 (2009)
17. Perc, M., Jordan, J.J., Rand, D.G., Wang, Z., Boccaletti, S., Szolnoki, A.: Statistical physics of human cooperation. *Phys. Rep.* 687, 1–51 (2017)
18. Lazer, D., Kennedy, R., King, G., Vespignani, A.: The parable of Google Flu: traps in big data analysis. *Science* 343(6176), 1203–1205 (2014)
19. Freeman, L.C.: Centrality in social networks: conceptual clarification. *Soc. Netw.* 1(3), 215–239 (1978)
20. Gershenson, C., Fernández, N.: Complexity and information: measuring emergence, self-organization, and homeostasis at multiple scales. *Complexity* 18(2), 29–44 (2012)
21. Hoel, E.P., Albantakis, L., Tononi, G.: Quantifying causal emergence shows that macro can beat micro. *PNAS* 110(49), 19790–19795 (2013)